

Optimizations Assignment 2

Alexander Sepelenco, Niall Sauvage

2023

Bad Spatial Locality

number of channels: 32..2048 (always powers of 2)

kernel order: 1, 3, 5, or 7

```
for ( c = 0; c < nchannels; c++ ) {  
    for ( x = 0; x < kernel_order; x++ ) {  
        for ( y = 0; y < kernel_order; y++ ) {  
            sum += image[w+x][h+y][c] * kernels[m][c][x][y];  
        }  
    }  
    output[m][w][h] = (float) sum;  
}
```

Better Spatial Locality

number of channels: 32..2048 (always powers of 2)

kernel order: 1, 3, 5, or 7

```
for ( x = 0; x < kernel_order; x++) {  
    for ( y = 0; y < kernel_order; y++ ) {  
        for ( c = 0; c < nchannels; c++ ) {  
            sum += image[w+x][h+y][c] * kernels[m][c][x][y];  
        }  
    }  
    output[m][w][h] = (float) sum;  
}
```

First Look At Loop unrolling

We can loop unroll our inner loop since nchannels can be big

```
for ( x = 0; x < kernel_order; x++) {  
    for ( y = 0; y < kernel_order; y++ ) {  
        for ( c = 0; c < nchannels; c+=4) { // power of 2  
            sum+=image[w+x][h+y][c]*kernels[m][c][x][y];  
            sum+=image[w+x][h+y][c+1]*kernels[m][c+1][x][y];  
            sum+=image[w+x][h+y][c+2]*kernels[m][c+2][x][y];  
            sum+=image[w+x][h+y][c+3]*kernels[m][c+3][x][y];  
        }  
    }  
    output[m][w][h] = (float) sum;  
}
```

More Loop unrolling

```
for ( x = 0; x < kernel_order; x++) {  
    for ( y = 0; y < kernel_order; y++ ) {  
        for ( c = 0; c < nchannels; c+=8) {  
            sum+=image[w+x][h+y][c]*kernels[m][c][x][y];  
            sum+=image[w+x][h+y][c+1]*kernels[m][c+1][x][y];  
            sum+=image[w+x][h+y][c+2]*kernels[m][c+2][x][y];  
            sum+=image[w+x][h+y][c+3]*kernels[m][c+3][x][y];  
            sum+=image[w+x][h+y][c+4]*kernels[m][c+4][x][y];  
            sum+=image[w+x][h+y][c+5]*kernels[m][c+5][x][y];  
            sum+=image[w+x][h+y][c+6]*kernels[m][c+6][x][y];  
            sum+=image[w+x][h+y][c+7]*kernels[m][c+7][x][y];  
        }  
    }  
    output[m][w][h] = (float) sum;  
}
```

Even More Loop unrolling

```
for ( x = 0; x < kernel_order; x++) {  
    for ( y = 0; y < kernel_order; y++ ) {  
        for ( c = 0; c < nchannels; c+=16) {  
            sum+=image[w+x][h+y][c]*kernels[m][c][x][y];  
            sum+=image[w+x][h+y][c+1]*kernels[m][c+1][x][y];  
            // 14 more times incrementing c index  
        }  
    }  
    output[m][w][h] = (float)sum;  
}
```

1D Arrays

Every array will be return in 1D notation

```
for ( x = 0; x < kernel_order; x++) {  
    for ( y = 0; y < kernel_order; y++) {  
        for ( c = 0; c < nchannels; c+=16) {  
            int image_offset = (w+x) * width_offset +  
                ((h+y) * nchannels) + c;  
            int kernel_total_offset =  
                m * kernel_offset +  
                x * kernel_order + y + c * ko2;  
            sum += image_1d[image_offset] *  
                kernel[kernel_total_offset];  
            sum += image_1d[image_offset + 1] *  
                kernel[kernel_total_offset + ko2 * 1];  
        }  
    }  
}
```

Kernel Not In Contiguous Memory

kernel isn't accessed in contiguous memory

```
sum += image_1d[image_offset] *  
      kernel[kernel_total_offset];  
sum += image_1d[image_offset + 1] *  
      kernel[kernel_total_offset + ko2 * 1];  
sum += image_1d[image_offset + 2] *  
      kernel[kernel_total_offset + ko2 * 2];  
sum += image_1d[image_offset + 3] *  
      kernel[kernel_total_offset + ko2 * 3];
```


Transpose

Transposing kernel makes access times faster for inner loop

```
for(int m = 0; m < nkernels; m++) {  
    for (int c = 0; c < nchannels; c++) {  
        for (int x = 0; x < kernel_order; x++) {  
            m_times_kernel_offset = m * kernel_offset;  
            c_times_ko2 = c * ko2;  
            x_times_kernel_order = x * kernel_order;  
            kernel_total_offset_precalc =  
                m_times_kernel_offset +  
                x_times_kernel_order + c_times_ko2;  
            t_kernel_total_offset_precalc =  
                m_times_kernel_offset +  
                x_times_kernel_order * nchannels + c;  
        }  
    }  
}
```

Loop Fusing/Loop Collapsing

Transposing kernel makes access times faster for inner loop

```
for(int n = 0; n < (nkernels*nchannels*kernel_order);  
    int m = n/(nchannels*kernel_order);  
    int c = (n%(nchannels*kernel_order))/kernel_order;  
    int x = (n%(nchannels*kernel_order))%kernel_order;  
    m_times_kernel_offset = m * kernel_offset;  
    c_times_ko2 = c * ko2;  
    x_times_kernel_order = x * kernel_order;  
    kernel_total_offset_precalc =  
        m_times_kernel_offset +  
        x_times_kernel_order + c_times_ko2;  
    t_kernel_total_offset_precalc =  
        m_times_kernel_offset +  
        x_times_kernel_order * nchannels + c;
```

Continued: Loop Fusing/Loop Collapsing

```
for(int y = 0; y < kernel_order; y++) {  
    t_kernel[t_kernel_total_offset_precalc] =  
        (double) kernel[kernel_total_offset_precalc+y];  
    t_kernel_total_offset_precalc += nchannels;  
}  
}
```

Parallel Transposing

Parallelise using omp

```
#pragma omp parallel for
for(int n = 0; n < (nkernels*nchannels*kernel_order);
    int m = n/(nchannels*kernel_order);
    int c = (n%(nchannels*kernel_order))/kernel_order;
    int x = (n%(nchannels*kernel_order))%kernel_order;
    ...
}
```

Parallel and Collapse Main Convolution

Parallelise using omp

```
#pragma omp parallel for
for (int n = 0; n < nkernels*width*height; n++) {
    int m = n/(width*height);
    int w = (n%(width*height))/height;
    int h = (n%(width*height))%height;
    ...
}
```

Align Kernel array

Aligning memory ensures better cache hits. This will also allow us to use vectors in the next slide.

```
double * t_kernel =  
    _mm_malloc(sizeof(double) *  
        nchannels * kernel_order *  
        kernel_order * nkernels, 16); // 16 byte aligned
```

Vectorisation

Unfortunately floats cause issues when summed up as the result may overflow, hence we are forced to work with

```
__m128d v4sum = _mm_setzero_pd();  
// 2 kernel_order for loops...  
for ( c = 0; c < nchannels; c+=2 ) {  
    __m128 v4image_1d =  
        _mm_loadu_ps(image_1d+image_offset+c);  
    __m128d v4image_1d_pd =  
        _mm_cvtps_pd(v4image_1d);  
  
    __m128d v4t_kernel_pd =  
        _mm_load_pd(t_kernel+kernel_total_offset+c);  
  
    __m128d product =  
        _mm_mul_pd(v4image_1d_pd, v4t_kernel_pd);  
    v4sum = _mm_add_pd(v4sum, product);  
}
```

Continuation: Vectorisation

```
c+=2;
v4image_1d =
    _mm_shuffle_ps(v4image_1d, v4image_1d,
        _MM_SHUFFLE(0, 0, 3, 2));
v4image_1d_pd = _mm_cvtps_pd(v4image_1d);

v4t_kernel_pd =
    _mm_load_pd(t_kernel+kernel_total_offset+c);

product = _mm_mul_pd(v4image_1d_pd, v4t_kernel_pd);
v4sum = _mm_add_pd(v4sum, product);
// repeat inner loop 12 more times
}
```


Final Step in Vectorisation

We must finally sum up, after finishing the kernel convolution

We repeat for every nkernels again.

Sum up using `hadd_pd`, which adds two double vectors together

Extract lower double using `cvtsd` function

```
v4sum = _mm_hadd_pd(v4sum, v4sum);  
output1d[m * width_times_height + w * height + h] =  
    (float) _mm_cvtsd_f64(v4sum);
```

Source Code

```
/* the fast version of matmul written by the student */
void student_conv(float *** restrict image, int16_t **** restrict kernels, float *** restrict output,
    int width, int height, int nchannels, int nkernels,
    int kernel_order)
{
    //int h, w, x, y, c, m;
    int y, x, c;
    int width_times_height = width * height;
    float * output1d = **output;
    float * image_1d = **image;
    int16_t * kernel = ***kernels;
    int ko2 = kernel_order * kernel_order;
    int width_offset = (height*kernel_order) * nchannels;
    int kernel_offset = nchannels * ko2;
    //int kernel_offset_increase = ko2 << 3;
    int image_offset_precalc;
    int kernel_total_offset_precalc;
    int image_offset;
    //double sum;
    int kernel_total_offset;
    double * t_kernel = _mm_malloc(sizeof(double) * nchannels * kernel_order * kernel_order * nkernels, 16);

    int m_times_kernel_offset;
    int c_times_ko2;
    int x_times_kernel_order;
    int t_kernel_total_offset_precalc;

    // VERY BAD TRANSPOSE ALGORITHM
    // USE AT OWN RISK
    // MIGHT CAUSE YOUR CAT TO RUN AWAY
    #pragma omp parallel for // maybe this will numb the pain
    for(int n = 0; n < (nkernels*nchannels*kernel_order); n++){
        int m = n/(nchannels*kernel_order);
        int c = (n%(nchannels*kernel_order))/kernel_order;
        int x = (n%(nchannels*kernel_order))%kernel_order;
        m_times_kernel_offset = m * kernel_offset;
        c_times_ko2 = c * ko2;
        x_times_kernel_order = x * kernel_order;
        kernel_total_offset_precalc = m_times_kernel_offset + x_times_kernel_order + c_times_ko2;
        t_kernel_total_offset_precalc = m_times_kernel_offset + x_times_kernel_order * nchannels + c;
        for(int y = 0; y < kernel_order; y++){
            t_kernel[t_kernel_total_offset_precalc] = (double) kernel[kernel_total_offset_precalc+y];
            t_kernel_total_offset_precalc += nchannels;
        }
    }
}
```

Continued: Source Code

```
#pragma omp parallel for
for (int n = 0; n < kernels*width*height; n++) {
    int m = n/(width*height);
    int w = (n%(width*height))/height;
    int h = (n%(width*height))%height;
    m_times_kernel_offset = m * kernel_offset;
    _m128d v4sum = _mm_setzero_pd();
    for (x = 0; x < kernel_order; x++) {
        image_offset_precalc = (w*x) * width_offset;
        kernel_total_offset_precalc = m_times_kernel_offset + x * kernel_order * nchannels;
        for (y = 0; y < kernel_order; y++) {
            image_offset = image_offset_precalc + (h*y) * nchannels;
            kernel_total_offset = kernel_total_offset_precalc + y * nchannels;
            for (c = 0; c < nchannels; c+=2) {
                // t_kernel is aligned hence the load instead of loads
                // cvtps convert float vector to double vector 4 -> 2 (lower) lanes
                // shuffle by moving 2 high values to take 2 low values spot
                // mul two double vectors, add them, and sum them up
                _m128d v4image_id = _mm_loadu_ps(image_id+image_offset+c);
                _m128d v4image_id_pd = _mm_cvtps_pd(v4image_id);

                _m128d v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

                _m128d product = _mm_mul_pd(v4image_id_pd, v4t_kernel_pd);
                v4sum = _mm_add_pd(v4sum, product);

                c+=2;
                v4image_id = _mm_shuffle_ps(v4image_id, v4image_id, _MM_SHUFFLE(0, 0, 3, 2));
                v4image_id_pd = _mm_cvtps_pd(v4image_id);

                v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

                product = _mm_mul_pd(v4image_id_pd, v4t_kernel_pd);
                v4sum = _mm_add_pd(v4sum, product);

                c+=2;
                v4image_id = _mm_loadu_ps(image_id+image_offset+c);
                v4image_id_pd = _mm_cvtps_pd(v4image_id);

                v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

                product = _mm_mul_pd(v4image_id_pd, v4t_kernel_pd);
                v4sum = _mm_add_pd(v4sum, product);

                c+=2;
                v4image_id = _mm_shuffle_ps(v4image_id, v4image_id, _MM_SHUFFLE(0, 0, 3, 2));
                v4image_id_pd = _mm_cvtps_pd(v4image_id);

                v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

                product = _mm_mul_pd(v4image_id_pd, v4t_kernel_pd);
                v4sum = _mm_add_pd(v4sum, product);

                c+=2;
                v4image_id = _mm_loadu_ps(image_id+image_offset+c);
                v4image_id_pd = _mm_cvtps_pd(v4image_id);

                v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

                product = _mm_mul_pd(v4image_id_pd, v4t_kernel_pd);
                v4sum = _mm_add_pd(v4sum, product);

                c+=2;
                v4image_id = _mm_shuffle_ps(v4image_id, v4image_id, _MM_SHUFFLE(0, 0, 3, 2));
            }
        }
    }
}
```

Continued: Source Code

```
v4image_1d_pd = _mm_cvtps_pd(v4image_1d);

v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

product = _mm_mul_pd(v4image_1d_pd, v4t_kernel_pd);
v4sum = _mm_add_pd(v4sum, product);

c+=2;
v4image_1d = _mm_loadu_ps(image_1d+image_offset+c);
v4image_1d_pd = _mm_cvtps_pd(v4image_1d);

v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

product = _mm_mul_pd(v4image_1d_pd, v4t_kernel_pd);
v4sum = _mm_add_pd(v4sum, product);

c+=2;
v4image_1d = _mm_shuffle_ps(v4image_1d, v4image_1d, _MM_SHUFFLE(0, 0, 3, 2));
v4image_1d_pd = _mm_cvtps_pd(v4image_1d);

v4t_kernel_pd = _mm_load_pd(t_kernel+kernel_total_offset+c);

product = _mm_mul_pd(v4image_1d_pd, v4t_kernel_pd);
v4sum = _mm_add_pd(v4sum, product);
    }
}
v4sum = _mm_hadd_pd(v4sum, v4sum); // add the two lanes together and put in lower lane
output1d[m * width_times_height + w * height + h] = (float) _mm_cvttsd_f64(v4sum); // extract lower double
}
}
```

Speed Increase

With test conv 128 128 7 256 256, we can get to nearly 90x speed increase.

```
sepelena@stoker:~/Concurrency$ make run
Executing: conv
David conv time: 134442611 microseconds
Student conv time: 1497555 microseconds
The total speed up time was 89.77x and 132945056 microseconds less
COMMENT: sum of absolute differences (0.000000) within acceptable range (0.062500)
```

Image Size and Speed Increase

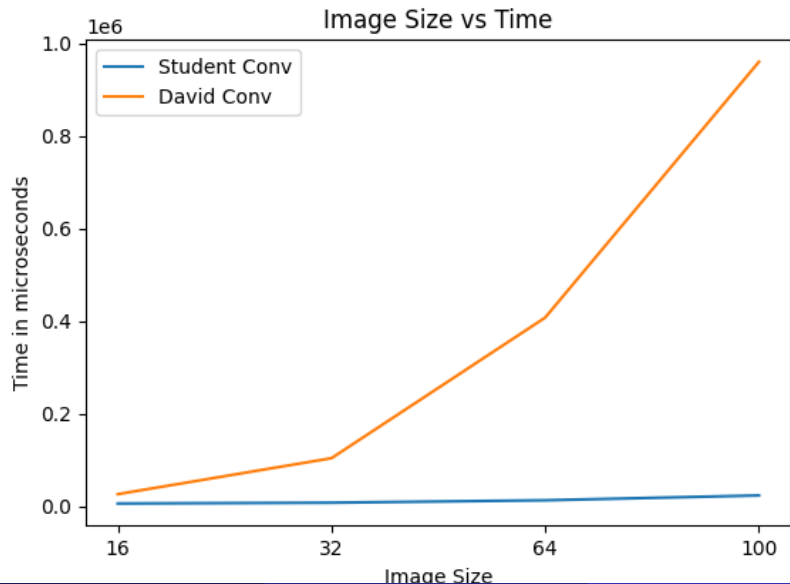



Image Size and Speed Increase



Graphs